

```

# Problem 9.7
library(ISLR)
# a
Auto$mpg=ifelse(Auto$mpg>median(Auto$mpg),1,0)
table(Auto$mpg)
# 0 1
# 196 196

# b
#install.packages("e1071")
library(e1071)
costs=data.frame(cost=seq(0.05,100,length.out = 15)) # tuning
grid for the cost parameter.
svm.tune=tune(svm,mpg~.,data=Auto,ranges=costs,kernel='linear') # 10-fold
cross validation.
svm.tune

# Parameter tuning of 'svm':
#
# - sampling method: 10-fold cross validation
#
# - best parameters:
#   cost
# 7.189286
#
# - best performance: 0.09918917
#
# - Detailed performance results:
#       cost      error dispersion
# 1  0.050000 0.10202964 0.03361002
# 2  7.189286 0.09918917 0.03261848
# 3 14.328571 0.10513783 0.03433267
# 4 21.467857 0.10826243 0.03387225
# 5 28.607143 0.11002209 0.03365457
# 6 35.746429 0.11084365 0.03392947
# 7 42.885714 0.11252353 0.03420288
# 8 50.025000 0.11337569 0.03417114
# 9 57.164286 0.11377659 0.03427571
# 10 64.303571 0.11408622 0.03453361
# 11 71.442857 0.11449827 0.03448589
# 12 78.582143 0.11487757 0.03433948
# 13 85.721429 0.11504715 0.03439838
# 14 92.860714 0.11566034 0.03452640
# 15 100.000000 0.11589444 0.03428601

# c
params=data.frame(cost=seq(0.05,100,length.out =
5),degree=seq(1,100,length.out = 5))
svm.poly=tune(svm,mpg~.,data=Auto,ranges=params,kernel='polynomial')

```

```

summary(svm.poly)
#
# Parameter tuning of 'svm':
#
# - sampling method: 10-fold cross validation
#
# - best parameters:
#   cost degree
# 100      1
#
# - best performance: 0.0974665

# - Detailed performance results:
#   cost degree      error dispersion
# 1    0.0500    1.00 0.31122882 0.04302732
# 2   25.0375    1.00 0.10085624 0.05523499
# 3   50.0250    1.00 0.09941058 0.05527120
# 4   75.0125    1.00 0.09868831 0.05421817
# 5  100.0000    1.00 0.09805010 0.05295848
# 6    0.0500   25.75 0.52345824 0.05422446
# 7   25.0375   25.75 0.52345824 0.05422446
# 8   50.0250   25.75 0.52345824 0.05422446
# 9   75.0125   25.75 0.52345824 0.05422446
# 10 100.0000   25.75 0.52345824 0.05422446
# 11    0.0500   50.50 0.52345824 0.05422446
# 12   25.0375   50.50 0.52345824 0.05422446
# 13   50.0250   50.50 0.52345824 0.05422446
# 14   75.0125   50.50 0.52345824 0.05422446
# 15 100.0000   50.50 0.52345824 0.05422446
# 16    0.0500   75.25 0.52345824 0.05422446
# 17   25.0375   75.25 0.52345824 0.05422446
# 18   50.0250   75.25 0.52345824 0.05422446
# 19   75.0125   75.25 0.52345824 0.05422446
# 20 100.0000   75.25 0.52345824 0.05422446
# 21    0.0500  100.00 0.52345824 0.05422446
# 22   25.0375  100.00 0.52345824 0.05422446
# 23   50.0250  100.00 0.52345824 0.05422446
# 24   75.0125  100.00 0.52345824 0.05422446
# 25 100.0000  100.00 0.52345824 0.05422446

params=data.frame(cost=seq(0.05,100,length.out =
5),gamma=seq(0.1,100,length.out = 5))
svm.radial=tune(svm,mpg~.,data=Auto,ranges=params,kernel='radial')
summary(svm.radial)
# Parameter tuning of 'svm':
#
# - sampling method: 10-fold cross validation
#
# - best parameters:

```

```

# cost gamma
# 25.0375 0.1
#
# # - best performance: 0.07497467
# cost gamma error dispersion
# 1 0.0500 0.100 0.07553539 0.020885918
# 2 25.0375 0.100 0.07070107 0.009828754
# 3 50.0250 0.100 0.07682088 0.009985928
# 4 75.0125 0.100 0.08041295 0.012046698
# 5 100.0000 0.100 0.08180043 0.013067447
# 6 0.0500 25.075 0.48294554 0.050646841
# 7 25.0375 25.075 0.24964179 0.002840959
# 8 50.0250 25.075 0.24964179 0.002840959
# 9 75.0125 25.075 0.24964179 0.002840959
# 10 100.0000 25.075 0.24964179 0.002840959
# 11 0.0500 50.050 0.48311419 0.050693482
# 12 25.0375 50.050 0.25124322 0.002629829
# 13 50.0250 50.050 0.25124322 0.002629829
# 14 75.0125 50.050 0.25124322 0.002629829
# 15 100.0000 50.050 0.25124322 0.002629829
# 16 0.0500 75.025 0.48313952 0.050694461
# 17 25.0375 75.025 0.25142340 0.002641727
# 18 50.0250 75.025 0.25142340 0.002641727
# 19 75.0125 75.025 0.25142340 0.002641727
# 20 100.0000 75.025 0.25142340 0.002641727
# 21 0.0500 100.000 0.48314265 0.050694752
# 22 25.0375 100.000 0.25144617 0.002644836
# 23 50.0250 100.000 0.25144617 0.002644836
# 24 75.0125 100.000 0.25144617 0.002644836
# 25 100.0000 100.000 0.25144617 0.002644836

```

```

# d
plot(svm.tune$performance[,c(1,2)],type='l')
# From the plot, we can see that while the cost of about 5 achieves the most
reduction in the miss classification error.

```

```


plot(svm.poly$performance[,c(2,1)],type='l')
# From the plot and from the summary, the best performance is achieved by
lowest degree of 1 and cost of 100.


```


```

plot(svm.radial$performance[,c(2,1)],type='l')
# From the plot and the summary, the best performance is achieved by a gamma of
0.1.

```







```

#Problem 10.2
#install.packages("knitr")
library(knitr)
# a
DissMatrix=data.frame(c(0,0.3,0.4,0.7),c(0.3,0,0.5,0.8),c(0.4,0.5,0,0.45),c(0.7,0.8,0.45,0),
colnames(DissMatrix)=c(paste('Col',1:4)))
kable(DissMatrix)
# | Col 1| Col 2| Col 3| Col 4|
# |-----:|-----:|-----:|-----:|
# | 0.0| 0.3| 0.40| 0.70|
# | 0.3| 0.0| 0.50| 0.80|
# | 0.4| 0.5| 0.00| 0.45|
# | 0.7| 0.8| 0.45| 0.00|

plot(hclust(dist(DissMatrix)),xlab='')
#The table gives the dissimilarity between classes, meaning that in order to
produce clusters we need only select pairs whose values are small. The
smallest value in the table is for the entries (1,2) and (2,1), followed by
the pair (3,4),(4,3).

# b
plot(hclust(dist(DissMatrix),method='single'),xlab='')

# c
#If we cut the first dendrogram obtain at the value of 0.7, then we obtain the
clusters (1,2) and (3,4). Cutting below this value produces 3 clusters since
the observations 3 and 4 are in their own cluster.

# d
#The clusters obtained here are (4) and (1,2,3).

# e
row.names(DissMatrix)=c(2,1,4,3)
plot(hclust(dist(DissMatrix)))
#A dendrogram is read bottom up, where the height indicates where clusters are
fused. Thus there is no horizontal meaning, the leafs are be swapped but they
still represent clusters that are fused at the same height.

# problem 10.7
USArrests.scaled=scale(USArrests)
correlation=as.dist(1-cor(t(USArrests.scaled)))
euclidean=dist(USArrests.scaled)^2

#If the quantities are approximately proportional then euclidean  $\approx$ 
K $\cdot$ correlation for a constant K.
summary(correlation/euclidean)

```

```
#   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#0.000086 0.069135 0.133943 0.234193 0.262589 4.887686
```

```
summary(correlation-0.1339*euclidean)
```

```
#   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
# -4.569956 -0.459715  0.000393 -0.059043  0.557616  1.808901
```

#If  $K=0.1339$  then they are approximately equal, different only 0.05 on average.

```
# Problem 10.10
```

```
# a
```

```
set.seed(42)
```

```
data= matrix(sapply(1:3,function(x){ rnorm(20*50,mean = 10*sqrt(x))
}),ncol=50) # 20 obs. in each class with 50 features.
```

```
class=unlist(lapply(1:3,function(x){rep(x,20)}))
```

```
# b
```

```
pr.out=prcomp(data)
```

```
plot(pr.out$x[,c(1,2)],col=class)
```

```
# c
```

```
set.seed(1)
```

```
kmeans.out=kmeans(data,3)
```

```
table(kmeans.out$cluster)
```

```
#  1  2  3
```

```
# 20 19 21
```

```
table(class)
```

```
# class
```

```
# 1  2  3
```

```
# 20 20 20
```

```
plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```

# Comparing to the graph from 10b, we can see that there is only one observation that is miss classified.

```
# d
```

```
set.seed(1)
```

```
kmeans.out=kmeans(data,2)
```

```
table(kmeans.out$cluster)
```

```
#1  2
```

```
#24 36
```

```
table(class)
```

```
# class
```

```
# 1  2  3
```

```
# 20 20 20
```



```
plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
#K-means seem to find a single cluster that is the same as before.
# We can see that k-mean separated the green middle section from the plot
10.d2, and put all the point the right of about 0.5 (PC1) to one cluster, and
left points to one cluster.
```

```
# e
set.seed(1)
kmeans.out=kmeans(data,4)
table(kmeans.out$cluster)
# 1  2  3  4
#19 10 17 14
```

```
plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```

#However, by examining the plot we can see that it again find the original green cluster with some overlap between it and the remaining ones. Overlap between clusters in the two principal components is also clear, #as should be expected since they may be close in the remaining dimensions.


```
# f
set.seed(1)
kmeans.out=kmeans(pr.out$x[,c(1,2)],3)
table(kmeans.out$cluster)
# 1  2  3
#20  8 32
plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
# This clustering seems to separated the plot where a clear path can be seen.
```

```
# g
set.seed(1)
kmeans.out=kmeans(scale(data,center = T,scale = T),3)
table(kmeans.out$cluster)
#1  2  3
#32 14 14
plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```

#There is significant overlap in the first two clusters, and the algorithm performs poorly compare to b.




Cluster Dendrogram



hclust (\*, "complete")

Cluster Dendrogram



hclust (\*, "single")

Cluster Dendrogram

